

Clustering: k-means

2024-09-05

- M1 MIDS/MFA
- [Université Paris Cité](#)
- Année 2024-2025
- [Course Homepage](#)

- [Moodle](#)



! Objectives

Setup

```
stopifnot(
  require(DT),
  require(skimr),
  require(GGally),
  require(patchwork),
  require(ggforce),
  require(glue),
  require(ggfortify),
  require(ggvoronoi),
  require(magrittr),
  require(broom),
  require(tidyverse)
)

tidymodels::tidymodels_prefer(quiet = TRUE)

old_theme <- theme_set(
  theme_minimal(base_size=9,
    base_family = "Helvetica")
)

knitr::opts_chunk$set(
  message = FALSE,
  warning = FALSE,
  comment=NA,
  prompt=FALSE,
  cache=FALSE,
  echo=TRUE,
  results='asis'
)
```

```
gc <- options(ggplot2.discrete.colour="viridis")
gc <- options(ggplot2.discrete.fill="viridis")
gc <- options(ggplot2.continuous.fill="viridis")
gc <- options(ggplot2.continuous.colour="viridis")
```

Foreword

This lab is dedicated to the *k-means* clustering method. In words, *k-means* takes as input a collection of points in \mathbb{R}^d (a numerical dataset) and a positive integer k . It returns a collection of k points (the *centers*) from \mathbb{R}^d . The centers define a Voronoï *tesselation/partition/diagram* of \mathbb{R}^d . The Voronoï *cells* define a clustering of the original dataset.

Voronoi tessellation/partition/diagram

[Wikipedia on Voronoï diagrams](#)

In the next chunk, we generate a Voronoï diagram on \mathbb{R}^2 with 100 cells defined from 100 random points drawn from the uniform distribution on a square. Function `stat_voronoi()` comes from `ggvoronoi`

```
set.seed(45056)
x <- sample(1:200,100)
y <- sample(1:200,100)
points <- tibble(x,
                 y,
                 distance = sqrt((x-100)^2 + (y-100)^2))

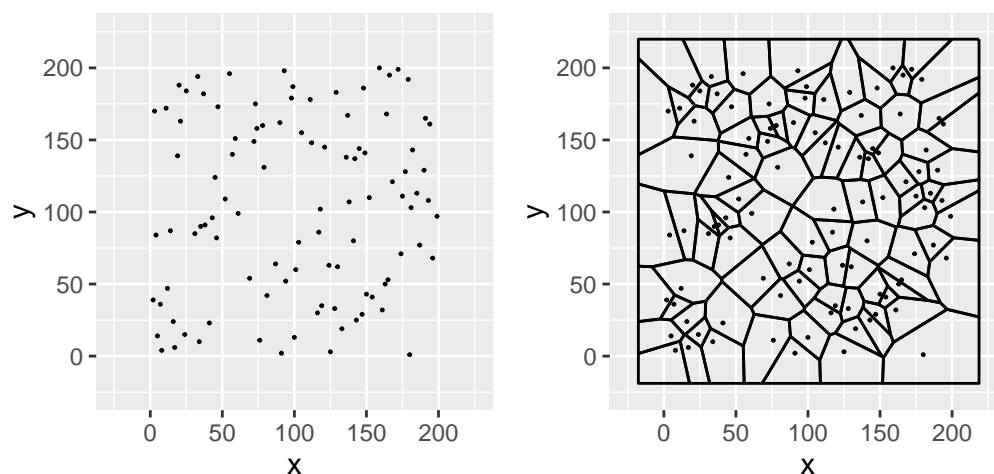
p <- ggplot(points) +
  aes(x=x, y=y) +
  geom_point(size=.2) +
  coord_fixed() +
  xlim(c(-25, 225)) +
  ylim(c(-25, 225))

p + (p + stat_voronoi(geom="path")) +
  patchwork::plot_annotation(
    title="Voronoi tessellation",
    subtitle = "Left: 100 random points\nRight: Voronoï diagram")
```

Voronoi tessellation

Left: 100 random points

Right: Voronoï diagram



- Two adjacent Voronoï cells are separated by a (possibly semi-infinite) line segment
- Let the so-called *centers* be denoted by c_1, \dots, c_n . They form the *codebook* \mathcal{C} .
- The Voronoï cell with *center* c_i is defined by

$$\left\{ x : x \in \mathbb{R}^d, \quad \|x - c_i\|_2 = \min_{j \leq n} \|x - c_j\|_2 \right\}$$

- The center of a Voronoï cell is usually not its barycenter

i *k*-means objective function

The *k*-means algorithm aims at building a *codebook* \mathcal{C} that minimizes

$$\mathcal{C} \mapsto \sum_{i=1}^n \min_{c \in \mathcal{C}} \|X_i - c\|_2^2$$

over all codebooks with given cardinality

If $c \in \mathcal{C}$ is the closest centroid to $X \in \mathbb{R}^p$,

$$\|c - X\|^2$$

is the *quantization/reconstruction error* suffered when using codebook \mathcal{C} to approximate X

⚠ If there are no restrictions on the dimension of the input space, on the number of centroids, or on sample size, computing an optimal codebook is a **NP** -hard problem

🔥 `kmeans()` is a wrapper for a collection of Algorithms that look like the Lloyd algorithm

Initialize by sampling from the data k-Mean++ try to take them as separated as possible.

Iterate the two phases until ?

👉 No guarantee to converge to a global optimum!

Proceed by trial and error.

Repeat and keep the best result.

Iris data

Run `kmeans()` on the projection of the Iris dataset on the Sepal plane. We look for a partition into three cells.

```
data(iris)

kms <- iris %>%
  select(starts_with("Petal")) %>%
  kmeans(3)
```

The result is an object of class `kmeans`. The class is equipped with `broom` methods.

Summarizing a clustering

i Question

Check the structure of objects of class `kmeans` and use `broom::tidy()` to get a summary. Compare with `summary()`

```
df_centers <- select(iris, starts_with("Petal")) |>
  kmeans(centers = 3) |>
  broom::tidy()

df_centers |>
  mutate(across(where(is.numeric), ~ signif(.x, digits=3) ))
```

```
# A tibble: 3 x 5
  Petal.Length Petal.Width size withinss cluster
    <dbl>         <dbl> <dbl>    <dbl> <fct>
1     4.29         1.36    54    14.2    1
2     1.46         0.246   50     2.02   2
3     5.63         2.05    46    15.2    3
```

Visualizing a clustering

i Question

Use `broom::augment()` and `broom::tidy()` to prepare two dataframes that will allow you to overlay a scatterplot of the dataset and a Voronoï diagram defined by the centers output by `kmeans()`.

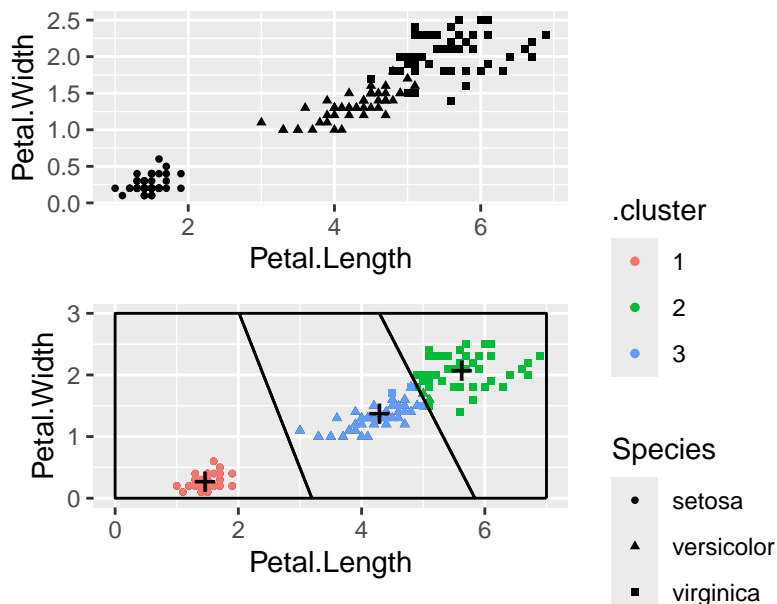
Compare the result with `plot()`

```
q <- kms %>%
  augment(iris) %>%
  ggplot() +
  aes(x=Petal.Length,
       y=Petal.Width
    ) +
  geom_point(aes(shape=Species), size=1, show.legend = F) +
  coord_fixed()

qq <- (q + geom_point(aes(shape=Species,
                          colour=.cluster),
                      size=1))+
  stat_voronoi(data = df_centers, #<<
              geom="path",
              outline=data.frame(x=c(0, 7, 7, 0),
                                y=c(0, 0, 3, 3))
            ) +
  geom_point(data = df_centers, #<<
            colour = "black",
            shape="+",
            size=5)

q / qq +
  plot_annotation(title = "Kmeans over Iris dataset, k=3")
```

Kmeans over Iris dataset, k=3



```
geom_sugar <- list(
  stat_voronoi(data = df_centers,
              geom="path",
              alpha=.5,
              outline = tribble(~x, ~y,
                                0., 0.,
                                7., 0.,
                                7., 3,
                                0., 3))
```

```

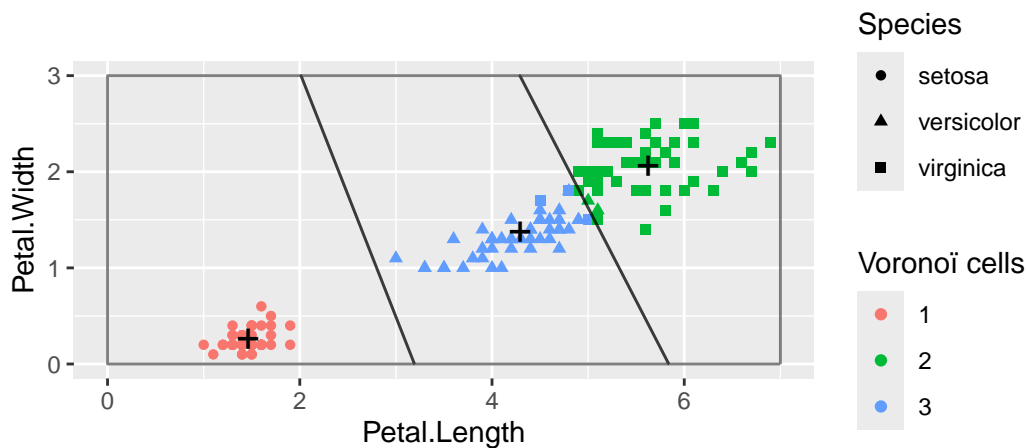
    ),
    geom_point(data = df_centers,
               colour = "black",
               shape="+",
               size=5),
    coord_fixed(),
    labs(col="Voronoi cells")
)

```

```

broom::augment(kms, iris) %>%
  ggplot(aes(x=Petal.Length, y=Petal.Width)) +
  geom_point(aes(shape=Species, color=.cluster)) +
  geom_sugar

```



i Question

Redo the same operations but choose the `Sepal.xxx` dimension.
Design a function to avoid repetitive coding.

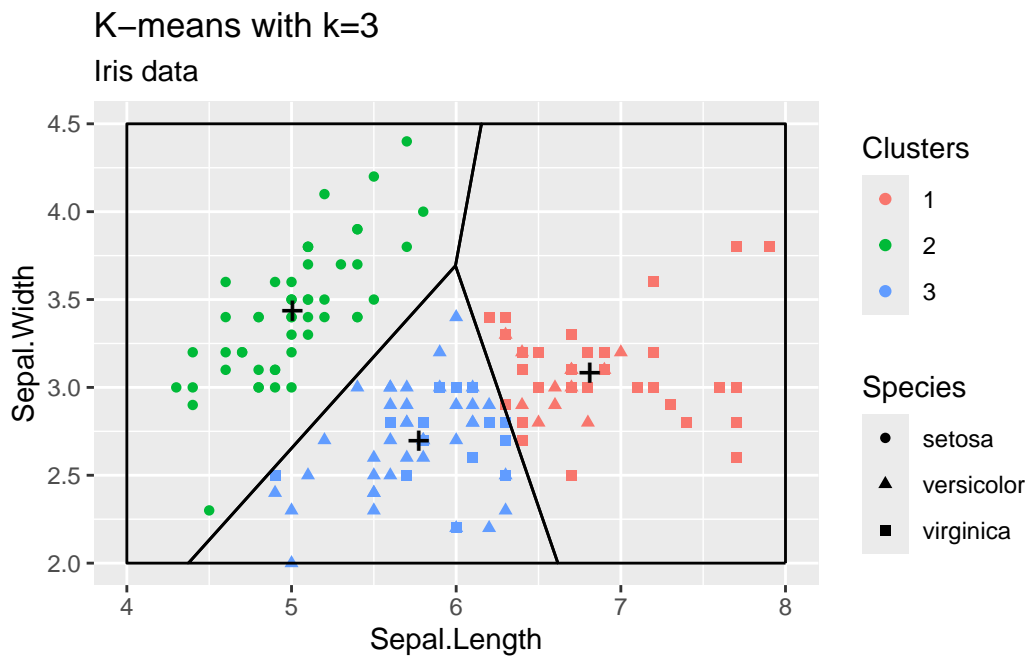
```

kms <- kmeans(select(iris, Sepal.Length, Sepal.Width), 3)

broom::augment(kms, iris) %>%
  ggplot() +
  geom_point(aes(x=Sepal.Length, y=Sepal.Width,
                 shape=Species, col=.cluster)) +
  geom_point(data=data.frame(kms$centers), #<<
             aes(x=Sepal.Length, y=Sepal.Width),
             shape='+',
             size=5) +
  stat_voronoi(data = as.data.frame(kms$centers), #<<
               aes(x=Sepal.Length, y=Sepal.Width),
               geom="path",
               outline=data.frame(x=c(4, 8, 8, 4),
                                 y=c(2, 2, 4.5, 4.5))) -> p

p +
  ggtitle("K-means with k=3", "Iris data") +
  labs(col="Clusters")

```



Playing with k

The number of cells/clusters may not be given a priori.

i Question

Perform kmeans clustering with $k = 2$. Use `glance`, `tidy`, `augment` to discuss the result.

```
df <- select(iris,
              starts_with("Sepal"))

kms <- kmeans(df, 2)

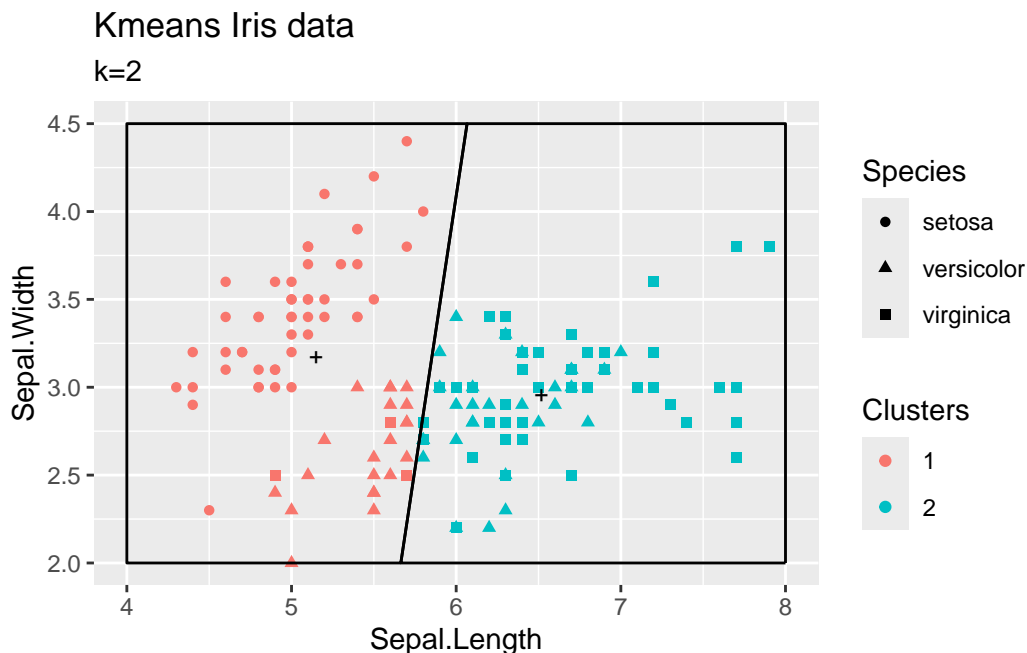
df_centers2 <- tidy(kms)

df_centers2

# A tibble: 2 x 5
  Sepal.Length Sepal.Width size withinss cluster
    <dbl>         <dbl> <int>    <dbl> <fct>
1     5.15         3.17    74     30.0  1
2     6.52         2.95    76     28.4  2

augment(kms, iris) %>%
  ggplot() +
  aes(x=Sepal.Length, y=Sepal.Width) +
  geom_point(aes(shape=Species, col=.cluster)) +
  geom_point(data=df_centers2, #<<
             shape='+',
             size=3) +
  stat_voronoi(data = df_centers2, #<<
               geom="path",
               outline=data.frame(x=c(4, 8, 8, 4),
                                   y=c(2, 2, 4.5, 4.5))) +
```

```
ggtitle(label="Kmeans Iris data",
        subtitle="k=2") +
labs(col="Clusters")
```



We can compare the spread between inner and outer sum of squares for clusterings with $k \in 2, 3$.

```
bind_rows(glance(kms),
          glance(kmeans(df,
                        centers=3,
                        nstart = 32L))) |>
mutate(k=c(2, 3))
```

```
# A tibble: 2 x 5
  totss tot.withinss betweenss iter    k
  <dbl>      <dbl>      <dbl> <int> <dbl>
1  130.         58.4        72.0     1     2
2  130.         37.1        93.4     3     3
```

i Question

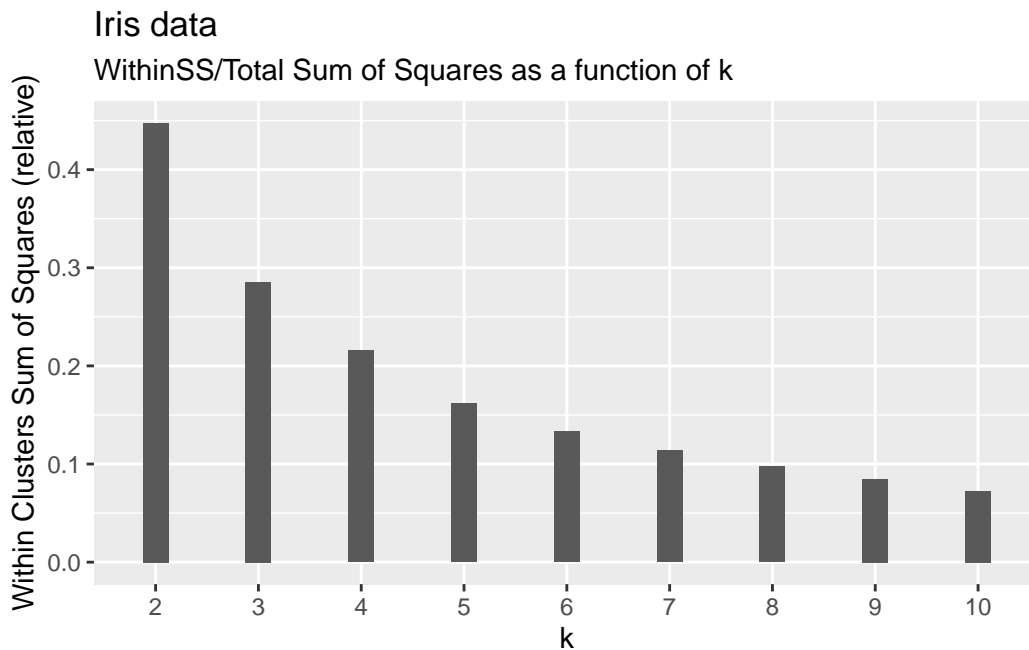
Perform k-means for $k = 2, \dots, 10$, plot within sum of squares as function of k . Comment.

```
tmp <-map_dfr(2:10, ~ glance(kmeans(df,
                                   centers=.,
                                   nstart = 32L))) %>%
rowid_to_column(var="k") %>%
mutate(k=k+1, across(where(is.numeric), ~ signif(.x, 3)))
```

```
tmp %>%
ggplot(aes(x=forcats::as_factor(k), y=tot.withinss/totss)) +
geom_col(width=.25) +
ggtitle("Iris data", "WithinSS/Total Sum of Squares as a function of k") +
xlab("k") +
ylab("Within Clusters Sum of Squares (relative)") +
```

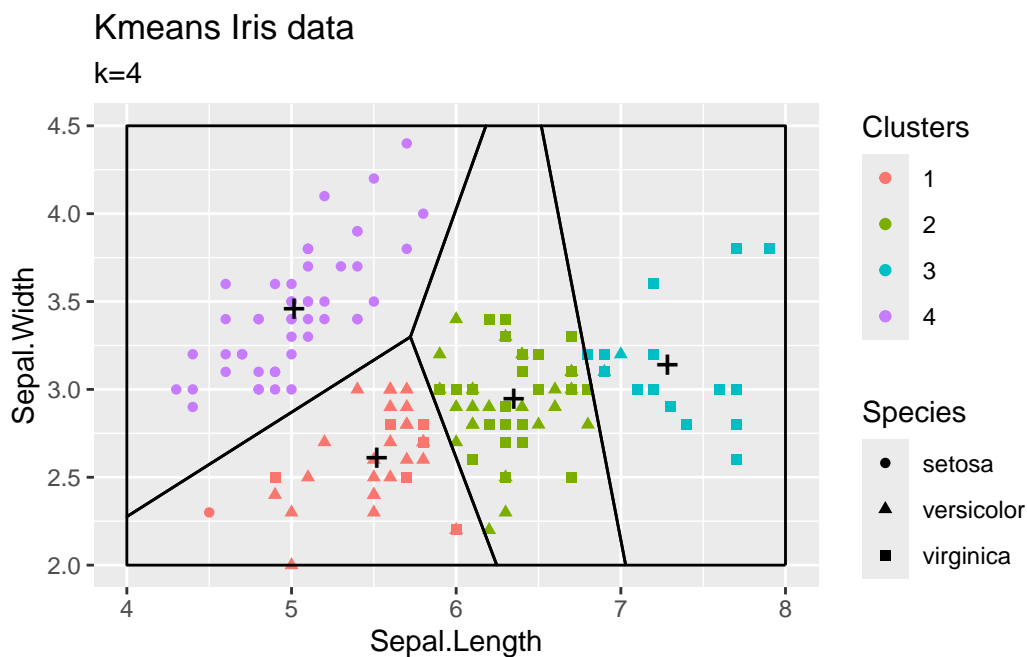


```
scale_x_discrete(breaks=as.character(2:10), labels=as.character(2:10))
```



```
kms <- kmeans(df, 4)
iris4 <- broom::augment(kms, iris)

ggplot(iris4) +
  geom_point(aes(x=Sepal.Length, y=Sepal.Width,
                 shape=Species, col=.cluster)) +
  geom_point(data=data.frame(kms$centers),          #<<
             aes(x=Sepal.Length, y=Sepal.Width),
             shape='+',
             size=5) +
  stat_voronoi(data = as.data.frame(kms$centers),  #<<
              aes(x=Sepal.Length, y=Sepal.Width),
              geom="path",
              outline=data.frame(x=c(4, 8, 8, 4), y=c(2, 2, 4.5, 4.5))) +
  ggtitle(label="Kmeans Iris data",
          subtitle="k=4") +
  labs(col="Clusters")
```



```

broom::tidy(kmeans(df, 4)) %>%
  knitr::kable(format = "markdown", digits = 2)

```

Sepal.Length	Sepal.Width	size	withinss	cluster
5.02	3.45	49	11.57	1
5.52	2.61	33	5.97	2
7.10	3.11	27	6.02	3
6.27	2.91	41	4.85	4

Lloyd's iterations

- i Initialize** Choose k centroids
Iterations: Two phases
- Movement** Assign each sample point to the closest centroid Assign each sample point to a class in the Voronoi partition defined by the centroids
- Update** For each class in the current Voronoi partition, update the centroid so as to minimize the Within Cluster Sum of Squared distances.

```
km <- list(centers=df[1:3, ]) # stupid initialization
```

```
sequence <- list()
```

```

for (i in 1:20) {
  km <- kmeans(df,
               km$centers,
               algorithm = "Lloyd",
               iter.max = 1)
  sequence[[length(sequence)+1]] <- force(km)
}

```

```

add_voronoi <- function(p, kmscenters, marker){
  p +
  geom_point(data=data.frame(kmscenters),

```

#<<

```

      mapping=aes(x=Sepal.Length, y=Sepal.Width),
      shape=marker,
      col="black",
      size=5) +
    stat_voronoi(data = as.data.frame(kmscenters), #<<
      aes(x=Sepal.Length,y=Sepal.Width),
      geom="path",
      outline=data.frame(x=c(4, 8, 8, 4),
        y=c(2, 2, 4.5, 4.5)))
  }

```

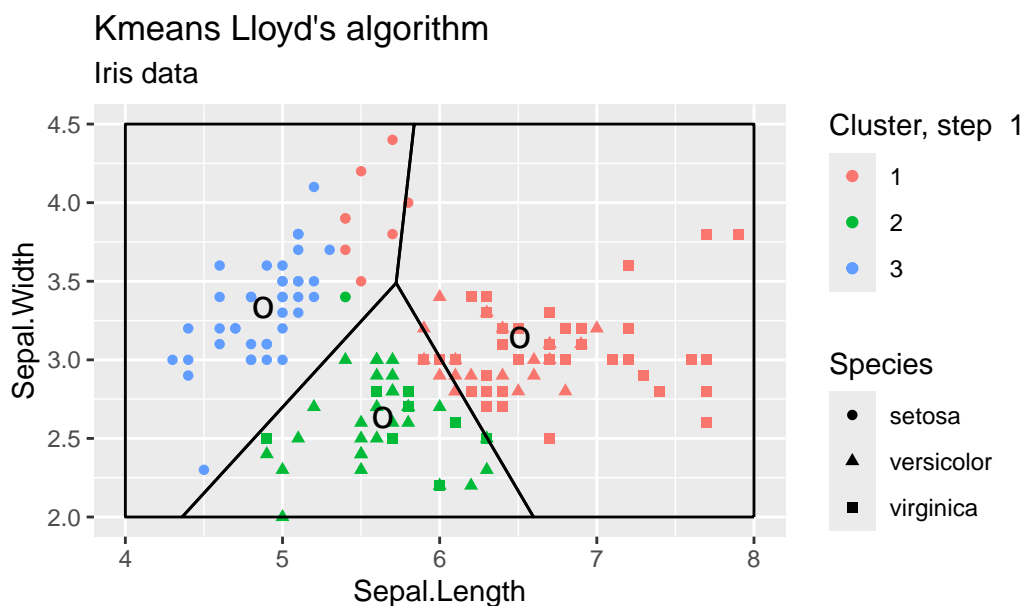
```
i <- 2
```

```

p <- broom::augment(sequence[[i]], iris) %>%
  ggplot() +
  coord_fixed(ratio=1) +
  geom_point(aes(x=Sepal.Length, y=Sepal.Width, shape=Species, col=.cluster)) +
  ggtitle("Kmeans Lloyd's algorithm", "Iris data")

p %>%
  add_voronoi(sequence[[i]]$centers, marker="o") + #<<
  labs(colour=paste("Cluster, step ", i- 1))

```

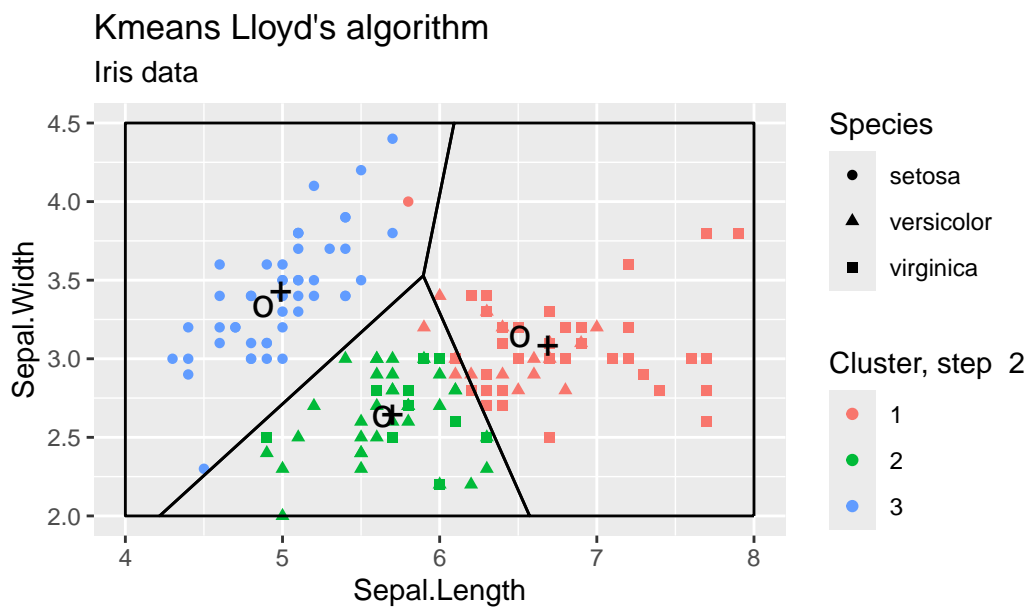


```
i <- 3
```

```

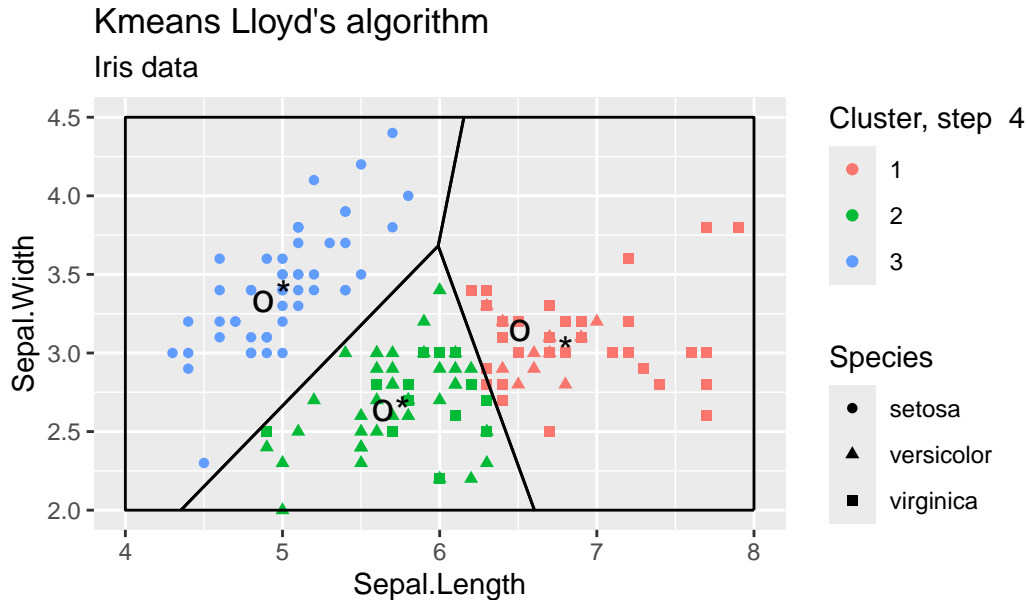
(p %>%
  broom::augment(sequence[[i]], iris)) %>%
  add_voronoi(sequence[[i]]$centers, marker='+') + #<<
  geom_point(data=data.frame(sequence[[2]]$centers), #<<
    mapping=aes(x=Sepal.Length, y=Sepal.Width),
    shape="o", col="black", size=5) +
  labs(colour=paste("Cluster, step ", i- 1))

```



```
i <- 5

(p %>%
  broom::augment(sequence[[i]], iris)) %>%
  add_voronoi(sequence[[i]]$centers, marker='*') + #<<
  geom_point(data=data.frame(sequence[[2]]$centers), #<<
    mapping=aes(x=Sepal.Length, y=Sepal.Width),
    shape="o", col="black",size=5) +
  labs(colour=paste("Cluster, step ", i- 1))
```



References

[Vignette k_means from tidyclust](#)