R language: a tour

2024-09-05

- M1 MIDS
- Université Paris Cité
- Année 2024-2025
- Course Homepage
- Moodle



Objectives

This workbook intends to walk you through basic aspects of the R language and programming environment.

Packages

Base R can do a lot. But the full power of R comes from a fast growing collection of packages.

Packages are first *installed* (that is downloaded from **cran** and copied somewhere on the hard drive), and if needed, *loaded* during a session.

- Installation can usually be performed using command install.packages(). In some circumstances, ad hoc installation commands (often from packages devtools) are needed
- Package pak offers an insteresting alternative to base R install.packages()
- Once a package has been installed/downloaded on your drive
 - if you want all objects exported by the package to be available in your session, you should load the package, using library() or require() (what's the difference?).
 Technically, this loads the NameSpace defined by the package.
 - if you just want to pick some objects exported from the package, you can use qualified names like package_name::object_name to access the object (function, dataset, ...).

For example, when we write

gapminder <- gapminder::gapminder</pre>

we assign dataframe/tibble gapminder from package gapminder to identifier "gapminder" in global environment .

Function p_load() from pacman (package manager) blends installation and loading: if the package named in the argument of p_load() is not installed (not among the installed.packages()), p_load() attempts to install the package. If installation is successful, the package is loaded.

```
if (! require(pak)){
  install.packages("pak")
}
to_be_loaded <- c("devtools",</pre>
                   "tidyverse",
                   "lobstr",
                   "ggforce",
                   "nycflights13",
                   "patchwork",
                   "glue",
                   "DT",
                   "kableExtra",
                   "viridis")
for (pck in to_be_loaded) {
  if (!require(pck, character.only = T)) {
    pak::pkg_install(pck, repos="http://cran.rstudio.com/")
    stopifnot(require(pck, character.only = T))
  }
}
```

Optional arguments

A very nice feature of R is that functions from base R as well as from packages have *optional* arguments with sensible *default* values. Look for example at documentation of require() using expression ?require.

Optional settings may concern individual functions or the collection of functions exported by some packages. In the next *chunk*, we reset the default color scales used by graphical functions from ggplot2.

```
opts <- options() # save old options

options(ggplot2.discrete.colour="viridis")
options(ggplot2.continuous.colour="viridis")</pre>
```

You shall not confuse *installing* (on your hard-drive) and *loading* (in session) a package.

i Question for Pythonistas

- In what is the analogue of install.packages()?
- In **?** what is the analogue of require()/library()?

Solution

In **\(\frac{1}{2} \)**, you can install a package pck using pip install pck or conda install pck (for example).

In **\(\big|**, the analogue of require(pck) could be

```
from pck import *
```

Note that in R, once a package in installed on the hard drive, you do not need to write something like

```
import pck
```

to be able to use objects exported by pck using qualified names (like pck.ze_object), you just need to use R qualified names:

```
pck::ze_object
```

Numerical (atomic) vectors

Numerical (atomic) vectors form the most primitive type of R.

Vector creation and assignment

The next three lines create three numerical atomic vectors.

In IDE Rstudio, have a look at the environment pane on the right before running the chunk, and after.

Use 1s() to investigate the *environment* before and after the execution of the three assignments.

```
ls()
x \leftarrow c(1, 2, 12)
y <- 5:7
z <- 10:1
x ; y ; z
ls()
```

Question

- What are the identifiers known in the global environment before execution of
- What are the identifiers known in the global environment after execution of lines 2-4?
- Which objects are attached to identifiers x, y, and z?

Solution

The chunks adds three identifiers x,y,z to the global environment. Identifiers are bound to R objects which turn out to be numerical vectors.

What does the next chunk?

```
ls()
w <- y
ls()
```

Solution

The chunk inserts a new identifier w in the global environment. This identifier is associated with the same object as y.

Question

- Is the content of object denoted by y copied to a new object bound to w?
- Interpret the result of w == y.
- Interpret the result of identical(w,y) (use help("identical") if needed).

```
identical(w,y)
```

Solution

Package lobstr lets us explore low-level aspects of R (and much more). Function

```
lobstr::obj_addr() returns the address of the object denoted by the argument.
lobstr::obj addr(w)
[1] "0x59e12ce35de8"
lobstr::obj_addr(y)
[1] "0x59e12ce35de8"
Now, if we modify either y or w
y < -y + 1
identical(y, w)
[1] FALSE
c(lobstr::obj_addr(w), lobstr::obj_addr(y))
[1] "0x59e12ce35de8" "0x59e12fe14c28"
The adress associated with y has changed!
```



The meaning of assignment in R differs from its countrepart in Python

Indexation, slicing, modification

Slicing a vector can be done in two ways:

- providing a vector of indices to be selected. Indices need not be consecutive
- providing a Boolean mask, that is a logical vector to select a set of positions

```
x \leftarrow c(1, 2, 12); y \leftarrow 5:7; z \leftarrow 10:1
```

Explain the next lines

```
z[1]  # slice of length 1
z[0]  # What did you expect?
z[x]  # slice of length ??? index error ?
z[y]
z[x %% 2]  # what happens with x[0] ?
z[0 == (x %% 2)]  # masking
z[c(2, 1, 1)]
```

Solution

- Indices start at 1 (not like in C, Java, or Python)
- $z \, [0]$ does not return an Error message. It returns an empty vector with the same basetype as x
- z[x] returns a vector made of z[x[1]], z[x[2]] and z[x[3]] == z[12]. Note again that z[12] does not raise an exception. It is simply not available (NA).
- x % 2 returns 1 0 0 as % stands for mod. z[x % 2] returns the same thing as z[1]
- c() stands for combine, or concatenate.

i Question

If the length of mask and and the length of the sliced vector do not coincide, what happens?

Solution

No error is signalled, the returned sequence is as long as the number of truthies in the mask.

Out of bound truthies show up as NA

```
z[rep(c(TRUE, FALSE), 6)]
[1] 10 8 6 4 2 NA
```

A scalar is just a vector of length 1!

```
class(z)
[1] "integer"

class(z[1])
[1] "integer"

class(z[c(2,1)])
[1] "integer"
```

Explain the next lines

```
y[2:3] \leftarrow z[2:3]

y == z[-10]

z[-11]
```

Solution

We can assign a slice of a vector to a slice of identical size of another vector. What is the result of z[-11], z[-c(11:7)]?

i Question

Explain the next line

```
z[-(1:5)]
```

Solution

We pick all positions in z but the ones in 1:5, that is 6, 7, 8, 9, 10

i Question

How would you select the last element from a vector (say z)?

Solution

```
z[length(z)]
```

[1] 1

1

Q is not **?** (reminder)!

i Question

Reverse the entries of a vector. Find two ways to do that.

Solution

```
z[seq(length(z), 1, by=-1)]
[1] 1 2 3 4 5 6 7 8 9 10
z[length(z):1]
[1] 1 2 3 4 5 6 7 8 9 10
rev(z) # the simplest way, once you know rev()
[1] 1 2 3 4 5 6 7 8 9 10
```

In statistics, machine learning, we are often faced with the task of building grid of regularly spaced elements (these elements can be numeric or not). R offers a collection of tools to

perform this. The most basic tool is rep().

• Question

- Repeat a vector 2 times
- Repeat each element of a vector twice

Solution w <- c(1, 7, 9) rep(w, 2) [1] 1 7 9 1 7 9 rep(w, rep(2, length(w))) [1] 1 1 7 7 9 9 Now, we can try something more fancy. rep(w, 1:3) [1] 1 7 7 9 9 9 What are the requirements on the second (times) argument?

Let us remove objects from the global environment.

```
rm(w, x, y ,z)
```

Numbers

So far, we told about numeric vectors. Numeric vectors are vectors of floating point numbers. R distinguishes several kinds of numbers.

- Integers
- Floating point numbers (double)

To check whether a vector is made of numeric or of integer, use is.numeric() or is.integer(). Use as.integer, as.numeric() to enforce type conversion.

```
Question
Explain the outcome of the next chunks
class(113L); class(113); class(113L + 113); class(2 * 113L); class(pi); as.intege
[1] "integer"
[1] "numeric"
[1] "numeric"
[1] "numeric"
[1] "numeric"
[1] 3
class(as.integer(113))
[1] "integer"
pi ; class(pi)
[1] 3.141593
[1] "numeric"
floor(pi) ; class(floor(pi)) # mind the floor
[1] 3
[1] "numeric"
```

Integer arithmetic

```
29L * 31L ; 899L %/% 32L ; 899L %% 30L

[1] 899

[1] 28

[1] 29
```

```
R integers are not the natural numbers from Mathematics
R numerics are not the real numbers from Mathematics
.Machine$double.eps
[1] 2.220446e-16
.Machine$double.xmax
[1] 1.797693e+308
.Machine$sizeof.longlong
[1] 8
u <- double(19L)
v <- numeric(5L)</pre>
w <- integer(7L)
lapply(list(u, v, w), typeof)
[[1]]
[1] "double"
[[2]]
[1] "double"
[[3]]
[1] "integer"
length(c(u, v, w))
[1] 31
typeof(c(u, v, w))
[1] "double"
```

R is (sometimes) able to make sensible use of Infinite.

```
log(0)
[1] -Inf
log(Inf)
[1] Inf
1/0
[1] Inf
0/0
[1] NaN
max(c( 0/0,1,10))
[1] NaN
max(c(NA,1,10))
[1] NA
max(c(-Inf,1,10))
[1] 10
```

```
is.finite(c(-Inf,1,10))
[1] FALSE TRUE TRUE
is.na(c(NA,1,10))
[1] TRUE FALSE FALSE
is.nan(c(NaN,1,10))
[1] TRUE FALSE FALSE
Computing with vectors
Summing, scalar multiplication
x <- 1:3
y < -9:7
sum(x); prod(x)
[1] 6
[1] 6
z <- cumsum(1:3)
w <- cumprod(3:5)</pre>
x + y
[1] 10 10 10
x + z
[1] 2 5 9
2 * w
[1] 6 24 120
2 + w
[1] 5 14 62
w / 2
[1] 1.5 6.0 30.0
  Question
    How would you compute a factorial?
  Solution
  n <- 10
  cumprod(1:n)
                     2
   [1]
             1
                             6
                                    24
                                           120
                                                   720
                                                           5040
                                                                  40320 362880
  [10] 3628800
```

Approximate $\sum_{n=1}^{\infty} 1/n^2$ within 10^{-3} ?

Solution

$$\sum_{n>N} \frac{1}{n^2} < \sum_{n>N} \frac{1}{n(n-1)} = \sum_{n>N} \left(\frac{1}{n-1} - \frac{1}{n} \right) = \frac{1}{N}$$

So we may pick N = 1000.

```
sum(x*y) # inner product
```

[1] 46

```
prod(1:5) # factorial(n) as prod(1:n)
```

[1] 120

```
N \leftarrow 1000L

sum(1/((1:N)^2)); pi^2/6 # grand truth
```

- [1] 1.643935
- [1] 1.644934

```
(pi^2/6 - sum(1/((1:N)^2))) < 1e-3
```

[1] TRUE

```
# N <- 999L
# (pi^2/6 - sum(1/((1:N)^2))) < 1e-3
```

i Question

How would you compute the inner product between two (atomic numeric) vectors?

Solution

Inner product between two vectors can be computed as a matrix product between a row vector and a column vector using %*%. Is this a good idea.

```
matrix(w, ncol=3) %*% matrix(y, nrow=3) == sum(w * y)
        [,1]
[1,] TRUE
```

- What we have called vectors so far are indeed atomic vectors.
 - Read Chapter on Vectors in R advanced Programming
 - Keep an eye on package vctrs for getting insights into the R vectors.

Numerical matrices

R offers a matrix class.

```
A <- matrix(1:50, nrow=5)
A
```

[2,]	2	7	12	17	22	27	32	37	42	47
[3,]	3	8	13	18	23	28	33	38	43	48
[4,]	4	9	14	19	24	29	34	39	44	49
[5,]	5	10	15	20	25	30	35	40	45	50

class(A)

[1] "matrix" "array"

i Question

From the evaluation of the preceding chunk, can you guess whether it is easier the traverse a matrix in row-first order or in column-first order?

Solution

Default traversal seems to proceed columnwise.

Creation, transposition and reshaping

A vector can be turned into a column matrix.

```
v <- as.matrix(1:5)
v</pre>
```

[,1]
[1,] 1
[2,] 2
[3,] 3
[4,] 4
[5,] 5

A matrix can be transposed

```
t(v) # transpose
```

```
[,1] [,2] [,3] [,4] [,5]
[1,] 1 2 3 4 5
cat(dim(v), ' ', dim(t(v)), '\n')
```

5 1 1 5

```
A <- matrix(1, nrow=5, ncol=2); A
```

```
[,1] [,2]
[1,] 1 1
[2,] 1 1
[3,] 1 1
[4,] 1 1
[5,] 1 1
```

lobstr::mem_used() allows us to keep track of the amount of memory used by
our R session. lobstr::obj_size() tells us the amount of memory used by the
representation of an object.

Comment the next chunk

```
m1 <-lobstr::mem_used()
A <- matrix(rnorm(100000L), nrow=1000L)
m2 <- lobstr::mem_used()
lobstr::obj_size(A)

800,216 B
B <- t(A)
lobstr::obj_size(B)

800,216 B

m3 <- lobstr::mem_used()
m2-m1; m3-m2

805,856 B
986,984 B</pre>
```

i e

Question

- Is there a difference between the next two assignments?
- How would you assign value to all entries of a matrix?

```
A <- matrix(rnorm(16), nrow=4)
A[] \leftarrow 0 ; A
     [,1] [,2] [,3] [,4]
              0
[1,]
                   0
[2,]
        0
              0
                   0
                         0
              0
                   0
                        0
[3,]
        0
                        0
[4,]
        0
              0
                   0
A < -0; A
[1] 0
```

Solution

There is!

The first assignment assigns 0 to every entry in A.

The second assignment binds 0 to name A

i Question What is the final shape of A? A <- matrix(1, nrow=5, ncol=2) A A[] <- 1:15 A</pre>

We can easily generate diagonal matrices and constant matrices.

```
diag(1, 3) # building identity matrix
     [,1] [,2] [,3]
[1,]
        1
              0
[2,]
        0
              1
                   0
[3,]
              0
matrix(0, 3, 3) # building null matrix
     [,1] [,2] [,3]
[1,]
        0
              0
[2,]
                   0
        0
              0
[3,]
        0
```

i Question

Is there any difference between the next two assignments?

```
B <- A[]
B ; A

[1] 0

[1] 0

lobstr::obj_addr(B) ; lobstr::obj_addr(A)

[1] "0x59e12e048168"

[1] "0x59e12f4ba6f8"

B <- A
```

Indexation, slicing, modification

Indexation consists in getting one item from a vector/list/matrix/array/dataframe.

Slicing and subsetting consists in picking a substructure:

- subsetting a vector returns a vector
- subsetting a list returns a list
- subsetting a matrix/array returns a matrix/array (beware of implicit simplifications and dimension dropping)
- subsetting a dataframe returns a dataframe or a vector (again, beware of implicit simplifications).

```
Explain the next results
```

```
A <- matrix(1, nrow=5, ncol=2)

dim(A[sample(5, 3), -1])
dim(A[sample(5, 3), 1])
length(A[sample(5, 3), 1])
is.vector(A[sample(5, 3), 1])
A[10:15]
A[60]
dim(A[])</pre>
```

NULL NULL [1] 3 [1] TRUE [1] 1 NA NA NA NA NA [1] NA [1] 5 2

i Question

How would you create a fresh copy of a matrix?

Solution

```
A <- matrix(rnorm(10), ncol=2L)
B <- matrix(0, nrow=5L, ncol=2L)

B[] <- A
all(B==A) ; identical(A, B) ; lobstr::obj_addrs(list(A, B))

[1] TRUE
[1] TRUE
[1] "0x59e12fe1cba8" "0x59e12fe1d3e8"</pre>
```

Computing with matrices

* versus %*% %*% stands for matrix multiplication. In order to use it, the two matrices should have conformant dimensions.

```
t(v) %*% A

[,1] [,2]
[1,] -1.243378 4.327579
```

There are a variety of reasonable products around. Some of them are available in R.

How would you compute the Hilbert-Schmidt inner product between two matrices?

$$\langle A,B\rangle_{\mathrm{HS}} = \mathrm{Trace}(A\times B^\top)$$

Solution

In R, trace() does not return the trace of a matrix! Function is used for debugging. Just remember that the trace of a matrix is the sum of its diagonal elements.

```
A <- matrix(runif(6), 2, 3)

B <- matrix(runif(6), 2, 3)

foo <- sum(diag(A %*% t(B)))

bar <- sum(A * B)

foo ; bar

[1] 2.093397

[1] 2.093397

Are you surprised?
```

Question

How can you invert a square (invertible) matrix?

Use solve(A) which is a shorthand for solve(A, diag(1, nrow(3))).

Logicals

- R has constants TRUE and FALSE.
- Numbers can be coerced to logicals.

Question

- Which numbers are truthies? falsies?
- What is the value (if any) of ! pi & TRUE ?
- What is the meaning of all()?
- What is the meaning of any()?
- Recall De Morgan's laws. Check them with R.
- Is | denoting an inclusive or an exclusive OR?

```
Solution
  w <- c(TRUE, FALSE, FALSE)
  sum(w)
  [1] 1
  any(w)
  [1] TRUE
  all(w)
  [1] FALSE
  ! w
  [1] FALSE TRUE TRUE
  TRUE & FALSE
  [1] FALSE
  TRUE | FALSE
  [1] TRUE
  TRUE | TRUE
  [1] TRUE
```

Handling three-valued logic

```
I Question

TRUE & (1> (0/0))
  (1> (0/0)) | TRUE
  (1> (0/0)) | FALSE
  TRUE || (1> (0/0))
  TRUE | (1> (0/0))
  TRUE | (1> (0/0))
  TRUE || stopifnot(4<3)
  # TRUE | stopifnot(4<3)
  # FALSE && stopifnot(4<3)
  # FALSE & stopifnot(4<3)</pre>
```

Solution TRUE & (1> (0/0)) [1] NA (1> (0/0)) | TRUE [1] TRUE (1> (0/0)) | FALSE [1] NA TRUE || (1> (0/0)) [1] TRUE TRUE | (1>(0/0))[1] TRUE TRUE || stopifnot(4<3) [1] TRUE # TRUE | stopifnot(4<3)</pre> FALSE && stopifnot(4<3)</pre> [1] FALSE # FALSE & stopifnot(4<3)</pre>

Question

What is the difference between logical operators | | and | ?

Solution

 $| \ |$ is lazy. It does not evaluate its second argument if the first one evaluates to TRUE. && is also lazy.

i 🖭

Remark: favor &, | over &&, | |.

all and any

Look at the definition of all and any.

i Question

- How would you check that a square matrix is symmetric?
- How would you check that a matrix is diagonal?

Solution

A square matrix is symmteric iff it is equal to its transpose. Recall that t(A) denotes the transpose of matrix A.

```
A <- matrix(rnorm(9), nrow=3, ncol=3) # a.s. non-symmetric all(A == t(A))
```

```
[1] FALSE

A <- A %*% t(A) # build a symmetric matrix, A + t(A) would work also
all(A == t(A))

[1] TRUE
A == t(A) returns a matrix a logical matrix, whose entries are all TRUE iff A is symmetric.
all() works for matrices as well as for vectors. This is sensible as matrices can be considered as vectors with some additional structure.</pre>
```

Lists

While an instance of an atomic vector contains objects of the same type/class, an instance of list may contain objects of widely different types.

```
Check an explain the output of the next chunk

p <- c(2, 7, 8)
q <- c("A", "B", "C")
x <- list(p, q)
x[2]
x
length(x)
rlang::is_vector(x)
rlang::is_atomic(x)
y <- c(p, q)
y
length(y)
rlang::is_atomic(y)
rlang::is_atomic(y)</pre>
```

```
[[1]]
[1] "A" "B" "C"

[[1]]
[1] 2 7 8

[[2]]
[1] "A" "B" "C"

[1] 2

[1] TRUE

[1] FALSE

[1] "2" "7" "8" "A" "B" "C"

[1] 6

[1] TRUE
```

[1] FALSE

- How would you build a list made of p, q, and x?
- What is x[2] made of?
- How does it compare with x[[2]]?

```
Solution
nl \leftarrow list(p=p, q=q, x=x)
$p
[1] 2 7 8
$q
[1] "A" "B" "C"
$x
$x[[1]]
[1] 2 7 8
$x[[2]]
[1] "A" "B" "C"
Note that we have defined a named list. Each = expression, binds the string on the
left-hand side to the object on the right-hand side. List elements can be extracted in
defferent ways.
names(nl)
[1] "p" "q" "x"
nl$q
[1] "A" "B" "C"
nl[["q"]]
[1] "A" "B" "C"
nl[[2]]
[1] "A" "B" "C"
```

Read and understand the next expressions.

```
is_atomic(p); is_atomic(p[2]); is_atomic(p[[2]])
is_list(q); is_atomic(q)
is_list(x); is_atomic(x); class(x)

class(x[2]); class(x[[2]])
length(x[2]); length(x[[2]])

identical(q, x[[2]]); identical(q, x[2])

obj_addr(q); obj_addr(x[[2]]); obj_addr(x[2])
ref(x)
obj_addrs(x)
identical(x[2],x[[2]])
```

i Functions is_atomic(), is_list(), ..., obj_addr() are from packages rlang and lobstr. See https://rlang.r-lib.org and https://lobstr.r-lib.org

```
Solution
[1] TRUE
[1] TRUE
[1] TRUE
[1] FALSE
[1] TRUE
[1] TRUE
[1] FALSE
[1] "list"
[1] "list"
[1] "character"
[1] 1
[1] 3
[1] TRUE
[1] FALSE
[1] "0x59e12baee2f8"
[1] "0x59e12baee2f8"
[1] "0x59e12de66978"
 [1:0x59e12c540ce8] <list>
 [2:0x59e128df5cc8] <dbl>
 [3:0x59e12baee2f8] <chr>
[1] "0x59e128df5cc8" "0x59e12baee2f8"
[1] FALSE
```

p and a are atomic vectors with different base types. They are not lists. A list like x is not an atomic vector.

Inspection of object addresses shows that when building x from objects p and q, objects bound to "p" and "q" are not copied.

Note that x[[2]] and x[2] are different objects, the former is one element list, the

```
second is an atomic vector.

ref(x[2])
  [1:0x59e12d7bb7f0] <list>
  [2:0x59e12baee2f8] <chr>
obj_addr(x[[2]])
  [1] "0x59e12baee2f8"
```

i Question

How would you replace "A" in x with "K"?

Solution

```
w <- c(2, 7, 8)
v <- c("A", "B", "C")
x <- list(w, v)
```

Read Chapter on Lists in R advanced Programming

Lookup tables (aka dictionaries) using named vectors

A lookup table maps strings to values. It can be implemented using named vectors. If we want to map: "seine" to "75", "loire" to "42", "rhone" to "69", "savoie" to "73" we can proceed in the following way:

```
codes <- c(75L, 42L, 69L, 73L)
names(codes) <- c("seine", "loire", "rhone", "savoie")

codes["rhone"]; codes["aube"]</pre>
```

rhone

69

<NA>

NA

Question

What is the class of codes?

Solution

```
names(codes)
[1] "seine" "loire" "rhone" "savoie"

class(codes); class(names(codes))
[1] "integer"
[1] "character"

is_atomic(codes); is_character(codes) ; is_integer(codes)
[1] TRUE
```

- [1] FALSE
- [1] TRUE

i Question

Capitalize the names used by codes

Package stringr offers a function str_to_title() that could be of interest.

```
Solution

names(codes) <- stringr::str_to_title(names(codes))

codes

Seine Loire Rhone Savoie

75 42 69 73
```

Factors

Factors exist in Base R. They play a very important role. Qualitative/Categorical variables are implemented as Factors.

Meta-package tidyverse offers a package dedicated to factor engineering: forcats.

```
yraw <- c("g1","g1","g2","g2","g2","g3")
print(yraw)

[1] "g1" "g1" "g2" "g2" "g2" "g3"

summary(yraw)

Length Class Mode
    6 character character

is.vector(yraw); is.atomic(yraw)</pre>
```

- [1] TRUE
- [1] TRUE

Question

yraw takes few values. It makes sense to make it a factor. How does it change the behavior of *generic* function summary?

Solution

Load the (celebrated) iris dataset, and inspect variable Species

```
data(iris)
species <- iris$Species
levels(species)
[1] "setosa" "versicolor" "virginica"
summary(species)</pre>
```

```
setosa versicolor virginica
50 50 50
```

i Question

We may want to collapse virginica and versicolor into a single level called versinica

forcats offer a function fct_collapse().

```
Solution

col_species <- forcats::fct_collapse(
   species,
   versinica = c("versicolor", "virginica")
)

summary(col_species)
   setosa versinica
   50   100</pre>
```

Factors are used to represent categorical variables.

Question

- Load the whiteside data from package MASS.
- Have a glimpse.
- Assign column Insul to y

```
whiteside <- MASS::whiteside # importing the whiteside data
# ?whiteside # what are the whiteside data about?

tibble::glimpse(whiteside)

Rows: 56
Columns: 3
$ Insul <fct> Before, Before, Before, Before, Before, Before, Before, Before, Femp <dbl> -0.8, -0.7, 0.4, 2.5, 2.9, 3.2, 3.6, 3.9, 4.2, 4.3, 5.4, 6.0, 6.~
$ Gas <dbl> 7.2, 6.9, 6.4, 6.0, 5.8, 5.8, 5.6, 4.7, 5.8, 5.2, 4.9, 4.9, 4.3,~
y <- whiteside$Insul # picking a factor column</pre>
```

- What is the class of y?
- Is y a vector
- Is y ordered? What does ordered mean here?
- What are the levels of y? How many levels has y?
- Can you slice y?
- What are the binary representations of the different levels of y?

```
Solution
is.factor(y); is.vector(y); is.ordered(y)
[1] TRUE
[1] FALSE
[1] FALSE
class(y)
[1] "factor"
levels(y)
[1] "Before" "After"
nlevels(y)
[1] 2
y[1:10] # yes we can
 [1] Before Before Before Before Before Before Before Before Before
Levels: Before After
pryr::bits(y[31]) # looks like the two levels are represented by integers
[1] "00000000 00000000 00000000 00000010"
```

i Question

Summarize factor y

```
Solution

summary(y) # counts

Before After
    26    30

table(y) # one-way contingency table

y
Before After
    26    30

table(y)/sum(table(y))*100 # one-way contingency table as percentages

y
Before After
46.42857 53.57143
```

```
table(y) %>%
  knitr::kable(col.names = c("Insulation", "Frequency"),
                caption = "Whiteside data") # Pb encoding sur machine windows
                           Table 1: Whiteside data
                            Insulation
                                      Frequency
                            Before
                                              26
                            After
                                              30
forcats::fct_count(y) %>%
  knitr::kable(col.names = c("Insulation", "Frequency"),
                caption = "Whiteside data")
                           Table 2: Whiteside data
                            Insulation
                                       Frequency
                            Before
                                              26
                            After
                                              30
```

Factors nuts and bolts

When coercing a vector (integer, character, ...) to a factor, use forcats::as_factor() rather than base R as.factor().



🕊 🖝 Useful function to make nice barplots when constructing barplots.

Recall that when you want to display counts for a univariate categorical sample, you use a barplot. It is often desirable to rank the levels according to the displayed statistics (usually a count).

This can be done in a seamless way using functions like forcats::fct_infreq().

```
forcats::fct_count(y, prop = TRUE)
```

```
# A tibble: 2 x 3
            n
  <fct> <int> <dbl>
1 Before
           26 0.464
2 After
           30 0.536
z <- sample(y, length(y), replace = TRUE) # permutation of whiteside$Insul
sort(forcats::fct_infreq(z))
                             # first level is most frequent one
```

- [1] After After After After After After After After After
- [11] After After After After After After After After
- [21] After After After After After After After After After After
- [31] After After Before Before Before Before Before Before Before
- [41] Before Befo
- [51] Before Before Before Before Before

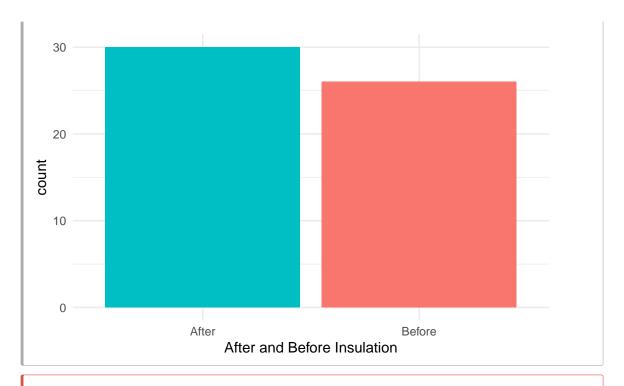
Levels: After Before

i Question

Make z ordered with level After preceding Before. Does ordering impact the behavior of forcats::fct_count()?

```
Solution
forcats::fct_count(z)
# A tibble: 2 x 2
 f
  <fct> <int>
1 Before
           24
2 After
           32
forcats::fct_count(factor(z, ordered=TRUE, levels=c("After", "Before")))
# A tibble: 2 x 2
 f
  <ord> <int>
1 After
          32
2 Before
           24
```

```
whiteside %>%
  ggplot2::ggplot() +
  ggplot2::aes(x=forcats::fct_infreq(Insul), fill=Insul) +
  ggplot2::geom_bar() +
  ggplot2::xlab("After and Before Insulation") +
  ggplot2::theme_minimal() +
  ggplot2::theme(legend.position="None")
```



Read Chapter on Factors in R for Data Science

Dataframes, tibbles and data.tables

A dataframe is a list of vectors with equal lengths. This is the way R represents and manipulates multivariate samples.

Any software geared at data science supports some kind of dataframe

- Python Pandas
- · Python Dask
- Spark
- ...

\$ Species

The iris dataset is the "Hello world!" of dataframes.

```
data(iris)
iris %>%
  glimpse()
```

```
Rows: 150

Columns: 5

$ Sepal.Length <dbl> 5.1, 4.9, 4.7, 4.6, 5.0, 5.4, 4.6, 5.0, 4.4, 4.9, 5.4, 4.~

$ Sepal.Width <dbl> 3.5, 3.0, 3.2, 3.1, 3.6, 3.9, 3.4, 3.4, 2.9, 3.1, 3.7, 3.~

$ Petal.Length <dbl> 1.4, 1.4, 1.3, 1.5, 1.4, 1.7, 1.4, 1.5, 1.4, 1.5, 1.5, 1.~

$ Petal.Width <dbl> 0.2, 0.2, 0.2, 0.2, 0.2, 0.4, 0.3, 0.2, 0.2, 0.1, 0.2, 0.~
```

<fct> setosa, setosa, setosa, setosa, setosa, setosa, setosa, s~

A matrix can be transformed into a data.frame

```
A <- matrix(rnorm(10), ncol=2)
data.frame(A)
```

```
X1 X2
1 0.06714445 -0.9791118
```

```
2 0.09486522 1.4178337
3 -0.88281225 0.9410125
4 -2.23246713 -0.3982120
5 -2.07713549 0.9225814
```

There are several flavors of dataframes in R: tibble and data.table are modern variants of data.frame.

```
t <- tibble::tibble(x=1:3, a=letters[11:13], d=Sys.Date() + 1:3)
head(t)</pre>
```

glimpse(t)

```
Rows: 3
Columns: 3
$ x <int> 1, 2, 3
$ a <chr> "k", "l", "m"
$ d <date> 2024-09-06, 2024-09-07, 2024-09-08
```

ref(t)

```
[1:0x59e12f98e048] <tibble[,3]>
x = [2:0x59e12c9d9230] <int>
a = [3:0x59e12fbb3348] <chr>
d = [4:0x59e12fa23f78] <date>
```

Read Chapter on data frames and tibbles in Advanced R

Question

Perform a random permutation of the columns of a data.frame/tibble.

¶ Function sample() from base R is very convenient

Solution

nycflights data

Wrestling with tables is part of the data scientist job. Out of the box data are often messy. In order to perform useful data analysis, we need tidy data. The notion of tidy data was elaborated during the last decade by experienced data scientists.

You may benefit from looking at the following online documents.

Tidy data in R for Data Science

Introduction to Table manipulation in R for Data Science in R.

More data of that kind is available following guidelines from https://github.com/hadley/nycflights13

In this exercise, you are advised to use functions from dplyr.

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges.

```
data <- nycflights13::flights
```

Question

- Have a glimpse at the data.
- What is the class of object data?
- What kind of object is data?

Hint: use class(), is.data.frame() tibble::is_tibble()

```
Solution
#| label: flight_glimpse
#| eval: true
data %>% glimpse()
Rows: 336,776
Columns: 19
              <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
$ year
$ month
              $ day
              $ dep_time
              <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
$ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
              <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
$ dep_delay
              <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849,~
$ arr_time
$ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851,~
              <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -1~
$ arr_delay
$ carrier
              <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
              <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
$ flight
```

```
<chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
$ tailnum
                 <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA",~
$ origin
                 <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD",~
$ dest
$ air_time
                 <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
                 <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
$ distance
$ hour
                 <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6
$ minute
                 <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
$ time_hour
                 <dttm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-01 0~
class(data)
[1] "tbl df"
                 "tbl"
                              "data.frame"
is.data.frame(data)
[1] TRUE
is tibble(data)
[1] TRUE
```

Extract the name and the type of each column.

```
Solution
colnames(data)
                             # name of columns
                       "month"
                                         "day"
 [1] "year"
                                                           "dep_time"
 [5] "sched_dep_time" "dep_delay"
                                         "arr_time"
                                                           "sched_arr_time"
 [9] "arr_delay"
                                                           "tailnum"
                       "carrier"
                                         "flight"
[13] "origin"
                       "dest"
                                                           "distance"
                                         "air_time"
[17] "hour"
                       "minute"
                                         "time_hour"
sapply(data, class)
                             # old school R, a dataframe is a list
$year
[1] "integer"
$month
[1] "integer"
$day
[1] "integer"
$dep_time
[1] "integer"
$sched_dep_time
[1] "integer"
$dep_delay
[1] "numeric"
$arr_time
[1] "integer"
```

```
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
                             # old school R, a dataframe is a list
lapply(data, class)
$year
[1] "integer"
$month
[1] "integer"
$day
[1] "integer"
$dep_time
[1] "integer"
$sched_dep_time
[1] "integer"
```

```
$dep_delay
[1] "numeric"
$arr_time
[1] "integer"
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
map(data, class)
                             # tidyverse way
$year
[1] "integer"
$month
[1] "integer"
$day
[1] "integer"
```

```
$dep_time
[1] "integer"
$sched_dep_time
[1] "integer"
$dep_delay
[1] "numeric"
$arr_time
[1] "integer"
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
```

Compute the mean of the numerical columns

Base R has plenty of functions that perform statistical computations on univariate samples. Look at the documentation of mean (just type ?mean). For a while, leave aside the optional

arguments.

In database parlance, we are performing aggregation

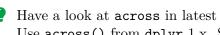
```
mean(data$dep_delay)
```

```
[1] NA
```

```
# mean(data[["dep_delay"]])
```

If we want the mean of all numerical columns, we need to project the data frame on numerical columns.

A verb of the summarize family can be useful.



Have a look at across in latest versions of dplyr() Use across() from dplyr 1.x. See Documentation

```
Solution
data %>%
  dplyr::select(where(is.numeric)) %>% # projecting on numerical columns
 purrr::map(mean)
                                         # applying the treatment to each column
$year
[1] 2013
$month
[1] 6.54851
$day
[1] 15.71079
$dep_time
[1] NA
$sched_dep_time
[1] 1344.255
$dep_delay
[1] NA
$arr_time
[1] NA
$sched_arr_time
[1] 1536.38
$arr_delay
[1] NA
$flight
[1] 1971.924
$air_time
```

```
[1] NA
$distance
[1] 1039.913
$hour
[1] 13.18025
$minute
[1] 26.2301
data %>%
  dplyr::select(where(is.numeric)) %>% # projecting on numerical columns
 dplyr::summarise(across(everything(), mean, na.rm=T))
# A tibble: 1 x 14
   year month
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
                                                 <dbl>
                                                          <dbl>
  <dbl> <dbl> <dbl>
                       <dbl>
                                       <dbl>
                                                                          <dbl>
1 2013 6.55 15.7
                       1349.
                                       1344.
                                                  12.6
                                                          1502.
                                                                          1536.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
data %>%
  dplyr::summarise(across(where(is.numeric), mean))
# A tibble: 1 x 14
   year month
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <dbl> <dbl> <dbl>
                       <dbl>
                                       <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                          <dbl>
1 2013 6.55 15.7
                          NA
                                       1344.
                                                    NA
                                                                          1536.
                                                             NA
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
data %>%
  dplyr::summarise(across(.cols=where(is.numeric), .fns=mean, na.rm=T))
# A tibble: 1 x 14
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   year month
  <dbl> <dbl> <dbl>
                       <dbl>
                                       <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                          <dbl>
1 2013 6.55 15.7
                       1349.
                                       1344.
                                                  12.6
                                                          1502.
                                                                          1536.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
```

If applied to a data.frame, summary(), produces a summary of each column. The summary depends on the column type. The output of summary is a shortened version the list of outputs obtained from applying summary to each column (lapply(data, summary)).

```
data %>%
  summary()
```

year	month	day	dep_time	sched_dep_time	
Min. :2013	Min. : 1.000	Min. : 1.00	Min. : 1	Min. : 106	
1st Qu.:2013	1st Qu.: 4.000	1st Qu.: 8.00	1st Qu.: 907	1st Qu.: 906	
Median :2013	Median : 7.000	Median :16.00	Median :1401	Median :1359	
Mean :2013	Mean : 6.549	Mean :15.71	Mean :1349	Mean :1344	
3rd Qu.:2013	3rd Qu.:10.000	3rd Qu.:23.00	3rd Qu.:1744	3rd Qu.:1729	
Max. :2013	Max. :12.000	Max. :31.00	Max. :2400	Max. :2359	
			NA's :8255		

```
dep_delay
                    arr_time
                               sched_arr_time
                                               arr_delay
                                                  : -86.000
Min.
      : -43.00
                 Min. : 1
                               Min. : 1
                                             Min.
1st Qu.: -5.00
                 1st Qu.:1104
                               1st Qu.:1124 1st Qu.: -17.000
Median: -2.00
                 Median:1535
                               Median:1556
                                             Median : -5.000
Mean
      : 12.64
                 Mean
                       :1502
                               Mean
                                      :1536
                                             Mean
                                                        6.895
3rd Qu.: 11.00
                 3rd Qu.:1940
                               3rd Qu.:1945
                                             3rd Qu.: 14.000
Max.
      :1301.00
                 Max.
                        :2400
                               Max.
                                      :2359
                                             Max.
                                                    :1272.000
NA's
      :8255
                 NA's
                        :8713
                                             NA's
                                                    :9430
  carrier
                     flight
                                  tailnum
                                                     origin
Length:336776
                  Min. : 1
                                Length:336776
                                                  Length: 336776
                                Class :character
                  1st Qu.: 553
Mode :character
                  Median:1496
                                Mode :character
                                                  Mode :character
                  Mean
                       :1972
                  3rd Qu.:3465
                  Max.
                         :8500
    dest
                     air_time
                                    distance
                                                    hour
Length: 336776
                  Min. : 20.0
                                 Min.
                                      : 17
                                               Min.
                                                      : 1.00
                  1st Qu.: 82.0
                                 1st Qu.: 502
Class : character
                                               1st Qu.: 9.00
Mode :character
                  Median :129.0
                                 Median: 872
                                               Median :13.00
                       :150.7
                                      :1040
                  Mean
                                 Mean
                                               Mean
                                                    :13.18
                  3rd Qu.:192.0
                                 3rd Qu.:1389
                                               3rd Qu.:17.00
                  Max.
                         :695.0
                                 Max. :4983
                                               Max.
                                                      :23.00
                  NA's
                         :9430
   minute
                 time_hour
Min. : 0.00
               Min.
                      :2013-01-01 05:00:00.00
1st Qu.: 8.00
               1st Qu.:2013-04-04 13:00:00.00
Median :29.00
               Median :2013-07-03 10:00:00.00
      :26.23
                      :2013-07-03 05:22:54.64
Mean
               3rd Qu.:2013-10-01 07:00:00.00
3rd Qu.:44.00
      :59.00
                     :2013-12-31 23:00:00.00
Max.
               Max.
```

Handling NAs

We add now a few NAs to the data....

```
data2 <- data
data2$arr_time[1:10] <- NA</pre>
```

Houston, we have a problem!

i Question

How should we compute the column means now?

Solution data2 %>% dplyr::summarise(across(is.numeric, mean)) # A tibble: 1 x 14 year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time

It is time to look at optional arguments of function mean.

i Question

Decide to ignore NA and to compute the mean with the available data

Solution

```
data2 %>%
 dplyr::summarise(across(is.numeric, mean, na.rm=TRUE))
# A tibble: 1 x 14
               day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  year month
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                <dbl>
                                                         <dbl>
1 2013 6.55 15.7
                       1349.
                                      1344.
                                                 12.6
                                                         1502.
                                                                         1536.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
   distance <dbl>, hour <dbl>, minute <dbl>
```

i Question

It is possible to remove all rows that contain at least one NA. Show this leads to a different result.

Solution

```
data2 %>%
  drop na() %>%
 dplyr::summarise(across(is.numeric, mean, na.rm=FALSE))
# A tibble: 1 x 14
               day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  year month
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                         <dbl>
1 2013 6.56 15.7
                       1349.
                                      1340.
                                                 12.6
                                                                         1533.
                                                          1502.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
   distance <dbl>, hour <dbl>, minute <dbl>
```

Question

Compute the minimum, the median, the mean and the maximum of numerical columns

```
data2 %>%
 dplyr::select_if(is.numeric) %>%
 lapply(function(x) c(med=median(x, na.rm=TRUE),
                     avg=mean(x, na.rm=TRUE),
                     max=max(x, na.rm=TRUE))
$year
med avg max
2013 2013 2013
$month
    med avg max
7.00000 6.54851 12.00000
$day
    med avg max
16.00000 15.71079 31.00000
$dep_time
         avg max
   med
1401.00 1349.11 2400.00
$sched_dep_time
    {\tt med} \qquad {\tt avg} \qquad {\tt max}
1359.000 1344.255 2359.000
$dep_delay
      med
              avg
 -2.00000 12.63907 1301.00000
$arr_time
    med avg max
1536.000 1502.075 2400.000
$sched_arr_time
               max
   med avg
1556.00 1536.38 2359.00
$arr_delay
       med avg max
 -5.000000 6.895377 1272.000000
$flight
    med avg max
1496.000 1971.924 8500.000
$air_time
            avg max
129.0000 150.6865 695.0000
$distance
```

```
med
              avg
                       max
 872.000 1039.913 4983.000
$hour
     med
              avg
                       max
13.00000 13.18025 23.00000
$minute
    med
            avg
                    max
29.0000 26.2301 59.0000
  data2 %>%
  dplyr::summarise(across(is.numeric, c(median=median,mean=mean, max=max), na.rm=T))
# A tibble: 1 x 42
  year_median year_mean year_max month_median month_mean month_max day_median
        <dbl>
                  <dbl>
                           <int>
                                         <dbl>
                                                    <dbl>
                                                               <int>
                                                                          <dbl>
         2013
                   2013
                             2013
                                                     6.55
1
                                                                  12
# i 35 more variables: day_mean <dbl>, day_max <int>, dep_time_median <int>,
    dep time mean <dbl>, dep time max <int>, sched dep time median <dbl>,
#
    sched_dep_time_mean <dbl>, sched_dep_time_max <int>,
#
    dep_delay_median <dbl>, dep_delay_mean <dbl>, dep_delay_max <dbl>,
#
    arr_time_median <int>, arr_time_mean <dbl>, arr_time_max <int>,
#
    sched_arr_time_median <dbl>, sched_arr_time_mean <dbl>,
    sched_arr_time_max <int>, arr_delay_median <dbl>, arr_delay_mean <dbl>, ...
```

Obtain a *nicer* output!

Check with https://dplyr.tidyverse.org/reference/scoped.html?q=funs#arguments

```
data2 %>%
  dplyr::summarise(across(is.numeric,
                          list(median=median,
                               mean=mean,
                               max=max)
                          na.rm=TRUE))
# A tibble: 1 x 42
  year_median year_mean year_max month_median month_mean month_max day_median
        <dbl>
                  <dbl>
                           <int>
                                         <dbl>
                                                    <dbl>
                                                              <int>
                                                                          <dbl>
         2013
1
                   2013
                            2013
                                             7
                                                     6.55
                                                                 12
                                                                             16
# i 35 more variables: day_mean <dbl>, day_max <int>, dep_time_median <int>,
    dep_time_mean <dbl>, dep_time_max <int>, sched_dep_time_median <dbl>,
    sched_dep_time_mean <dbl>, sched_dep_time_max <int>,
#
#
    dep_delay_median <dbl>, dep_delay_mean <dbl>, dep_delay_max <dbl>,
#
    arr_time_median <int>, arr_time_mean <dbl>, arr_time_max <int>,
    sched_arr_time_median <dbl>, sched_arr_time_mean <dbl>,
#
    sched_arr_time_max <int>, arr_delay_median <dbl>, arr_delay_mean <dbl>, ...
```

-2

600

Question

Mimic summary on numeric columns

```
Solution
mysum <- data2 %>%
  dplyr::summarise(across(is.numeric,
                          list(median=median,
                                mean=mean,
                                max=max,
                                min=min,
                                sd=sd,
                                IQR=IQR) ,
                      na.rm=TRUE))
mysum
# A tibble: 1 x 84
  year_median year_mean year_max year_min year_sd year_IQR month_median
        <dbl>
                  <dbl>
                            <int>
                                     <int>
                                             <dbl>
                                                       <dbl>
                                                                    <dbl>
1
         2013
                   2013
                             2013
                                      2013
                                                 0
# i 77 more variables: month_mean <dbl>, month_max <int>, month_min <int>,
#
    month_sd <dbl>, month_IQR <dbl>, day_median <dbl>, day_mean <dbl>,
    day_max <int>, day_min <int>, day_sd <dbl>, day_IQR <dbl>,
#
#
    dep_time_median <int>, dep_time_mean <dbl>, dep_time_max <int>,
#
    dep_time_min <int>, dep_time_sd <dbl>, dep_time_IQR <dbl>,
#
    sched_dep_time_median <dbl>, sched_dep_time_mean <dbl>,
    sched_dep_time_max <int>, sched_dep_time_min <int>, ...
```

i Question

Compute a new itinerary column concatenating the origin and dest one. Have a look at Section Operate on a selection of variables

```
Solution
```

ORD

LGA

10 ORD-LGA

```
data %>%
  dplyr::mutate(itinerary=paste(dest, origin, sep="-")) %>%
 dplyr::select(itinerary, dest, origin, everything())
# A tibble: 336,776 x 20
   itinerary dest origin year month
                                          day dep_time sched_dep_time dep_delay
   <chr>
             <chr> <chr> <int> <int> <int>
                                                 <int>
                                                                            <dbl>
                                                                 <int>
1 IAH-EWR
                   EWR
             IAH
                            2013
                                      1
                                            1
                                                   517
                                                                   515
                                                                                2
 2 IAH-LGA
             IAH
                   LGA
                            2013
                                      1
                                            1
                                                   533
                                                                   529
                                                                                4
                                                                                2
3 MIA-JFK
             MIA
                   JFK
                            2013
                                      1
                                            1
                                                   542
                                                                   540
                                                                               -1
4 BQN-JFK
             BQN
                   JFK
                            2013
                                            1
                                      1
                                                   544
                                                                   545
                                            1
5 ATL-LGA
             ATL
                   LGA
                            2013
                                      1
                                                                               -6
                                                   554
                                                                   600
6 ORD-EWR
             ORD
                   EWR
                            2013
                                      1
                                            1
                                                                   558
                                                                               -4
                                                   554
7 FLL-EWR
                                            1
                                                                               -5
             FLL
                   EWR
                            2013
                                      1
                                                   555
                                                                   600
8 IAD-LGA
             IAD
                   LGA
                            2013
                                      1
                                            1
                                                   557
                                                                   600
                                                                               -3
9 MCO-JFK
             MCO
                    JFK
                                            1
                                                                               -3
                            2013
                                      1
                                                   557
                                                                   600
```

1

1

558

2013

```
# i 336,766 more rows
# i 11 more variables: arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
# carrier <chr>, flight <int>, tailnum <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>
```

Compute the coefficient of variation (ratio between the standard deviation and the mean) for each itinerary. Can you find several ways?

```
Solution
data %>%
  dplyr::mutate(itinerary=paste(dest, origin, sep="-")) %>%
 dplyr::select(itinerary, dest, origin, everything()) %>%
 dplyr::group_by(itinerary) %>%
  dplyr::summarise(coef_var=sd(air_time, na.rm=T)/mean(air_time, na.rm=T), .groups = "dr
  slice_sample(n=10)
# A tibble: 10 x 2
  itinerary coef_var
                <dbl>
   <chr>
1 SFO-JFK
               0.0485
2 MTJ-EWR
               0.0541
3 DEN-EWR
               0.0688
4 BQN-EWR
               0.0453
5 IAH-EWR
               0.0824
6 HNL-JFK
              0.0332
7 RDU-LGA
               0.0861
8 PHL-JFK
               0.178
9 GRR-EWR
               0.0889
10 PIT-LGA
               0.0801
```

Question

Compute for each flight the ratio between the distance and the air_time in different ways and compare the execution time (use Sys.time()).

Solution before <- Sys.time()

```
3 MYR-LGA 0.166
4 CAE-LGA 0.164
5 AVL-LGA 0.154
6 LEX-LGA 0.149
7 SBN-LGA 0.149
8 SBN-EWR 0.147
9 JAC-JFK 0.145
10 STL-JFK 0.145
# i 214 more rows

required_time <- Sys.time() - before required_time

Time difference of 0.2313318 secs
```

Which carrier suffers the most delay?

```
Solution
data %>%
 dplyr::select(carrier, arr_delay) %>%
 dplyr::filter(arr_delay > 0) %>%
 dplyr::group_by(carrier) %>%
 dplyr::summarise(ndelays= n()) %>%
# dplyr::arrange(desc(ndelays)) %>%
# head(3)
 dplyr::top_n(3, ndelays)
# A tibble: 3 x 2
 carrier ndelays
  <chr> <int>
           23609
1 B6
2 EV
          24484
3 UA
           22222
```

Puzzle

```
year <- 2012L
data %>%
  dplyr::select(year, dest, origin) %>%
 head()
# A tibble: 6 x 3
  year dest origin
  <int> <chr> <chr>
1 2013 IAH EWR
2 2013 IAH
             LGA
3 2013 MIA
            JFK
4 2013 BQN
             JFK
5 2013 ATL
             LGA
6 2013 ORD
             EWR
```

```
data %>%
  dplyr::filter(year==year) %>%
  dplyr::summarize(n())
# A tibble: 1 x 1
   `n()`
   <int>
1 336776
data %>%
  dplyr::filter(year==2012L) %>%
  dplyr::summarize(n())
# A tibble: 1 x 1
  `n()`
  <int>
1
      0
data %>%
  dplyr::filter(year==.env$year) %>%
  dplyr::summarize(n())
# A tibble: 1 x 1
  `n()`
  <int>
1
      0
data %>%
  dplyr::filter(year==.data$year) %>%
  dplyr::summarize(n())
# A tibble: 1 x 1
   `n()`
   <int>
1 336776
```

Can you explain what happens?

Solution

When dplyr::filter(year==year) does year refer to the column of data or to the variable in the global environment?

Flow control

R offers the usual flow control constructs:

- branching/alternative if (...) {...} else {...}
- iterations (while/for) while (...) {...} for (it in iterable) {...}
- function calling callable(...) (how do we pass arguments? how do we rely on defaults?)

If () then {} else {}

If expressions yes_expr and no_expr are complicated it makes sense to use the if (...) {...} else {...} construct

There is also a conditional statement with an optional else {}

```
#| eval: false
#| collapse: false
if (condition) {
    ...
} else {
    ...
}
```

i Question

Is there an elif construct in R?

Nope!

R also offers a switch

```
switch (object,
  case1 = {action1},
  case2 = {action2},
  ...
)
```

```
i There exists a selection function ifelse(test, yes_expr, no_expr).
ifelse(test, yes, no)
Note that ifelse(...) is vectorized.

x <- 1L:6L
y <- rep("odd", 6)
z <- rep("even", 6)

ifelse(x %% 2L, y, z)
[1] "odd" "even" "odd" "even" "odd" "even"

This is a vectorized function</pre>
```

Iterations for (it in iterable) {...}

Have a look at Iteration section in R for Data Science

i Question

Create a lower triangular matrix which represents the 5 first lines of the Pascal triangle.

Recall

$$\binom{n}{k} = \binom{n-1}{k-1} + \binom{n-1}{k}$$

```
Solution

T <- matrix(OL, nrow=6, ncol=6)
T[1,1] <- 1L

for (i in 2:ncol(T))
   T[i, 1:i] <- c(OL, T[i-1, 2:i-1]) + T[i-1, 1:i]

colnames(T) <- OL:5L
rownames(T) <- OL:5L

T

   0 1 2 3 4 5
0 1 0 0 0 0 0
1 1 1 0 0 0 0
2 1 2 1 0 0 0
3 1 3 3 1 0 0
4 1 4 6 4 1 0
5 1 5 10 10 5 1
```

Locate the smallest element in a numerical vector

```
Solution

v <- sample(1:100, 100)
v[1:10]

[1] 48 75 41 92 83  3 56  9 84 20

pmin <- 1

for (i in seq_along(v)) {
   if (v[i] < v[pmin]) {
      pmin <- i
    }
}

print(stringr::str_c('minimum is at ', pmin, ', it is equal to ', v[pmin]))

[1] "minimum is at 15, it is equal to 1"
There are some redundant braces {}</pre>
```

While (condition) {...}

Question

Find the location of the minimum in a vector **v**

Write a loop that checks whether vector **v** is non-decreasing.

```
result <- TRUE

for (i in 2:length(v))
  if (v[i] < v[i-1]) {
    result <- FALSE
    break
  }

if (result) {
  print("non-decreasing")
} else {
  print("not non-decreasing")
}</pre>
```

Functions

To define a function, whether named or not, you can use the function constructor.

```
foo <- function(arg1, arg2=def2) {
    # body
}</pre>
```

Write a function that checks whether vector **v** is non-decreasing.

```
Solution
is_non_decreasing <- function(v) {</pre>
  for (i in 2:length(v))
    if (v[i] < v[i-1]) {</pre>
      return(FALSE)
    }
  return(TRUE)
}
is_non_decreasing(v)
[1] FALSE
is_non_decreasing(1:10)
[1] TRUE
A function is an object like any other
is_non_decreasing
function(v) {
  for (i in 2:length(v))
    if (v[i] < v[i-1]) {
      return(FALSE)
    }
  return(TRUE)
}
<bytecode: 0x59e13632d0c8>
body(is_non_decreasing)
{
    for (i in 2:length(v)) if (v[i] < v[i-1]) {
        return(FALSE)
    }
    return(TRUE)
}
args(is_non_decreasing)
function (v)
NULL
```

i Question

Write a function with integer parameter n, that returns the Pascal Triangle with n+1 rows.

```
triangle_pascal <- function(n) {</pre>
 m <- n+1
 T \leftarrow matrix(c(rep(1, m), rep(0, m*(m-1))), nrow=m, ncol=m)
 for (i in 2:m)
   T[i, 2:i] \leftarrow T[i-1, 1:i-1] + T[i-1, 2:i]
 for (i in 1:(m-1))
   T[i, (i+1):m] <- NA
 colnames(T) <- 0:n</pre>
 rownames(T) <- 0:n
 T
}
print(triangle_pascal(10), na.print=" " )
  0 1 2 3 4 5 6 7 8 9 10
0
  1
1 1 1
2 1 2 1
3 1 3 3
           1
4 1 4 6 4 1
5 1 5 10 10 5 1
6 1 6 15 20 15 6 1
7 1 7 21 35 35 21
                      7
                          1
8 1 8 28 56 70 56 28
                          8 1
9 1 9 36 84 126 126 84 36 9 1
10 1 10 45 120 210 252 210 120 45 10 1
Sanity check: R provides us with function choose
n <- 5
map(0:n, ~ choose(., 0:.))
[[1]]
[1] 1
[[2]]
[1] 1 1
[[3]]
[1] 1 2 1
[[4]]
[1] 1 3 3 1
[[5]]
[1] 1 4 6 4 1
[[6]]
[1] 1 5 10 10 5 1
```

```
t10 <- triangle_pascal(10)

for (n in 0:10)
  for (p in 0:n)
    stopifnot(t10[as.character(n), as.character(p)] == choose(n, p))</pre>
```

Question

How would you generate a Fibonacci sequence of length n?

Recall the Fibonacci sequence is defined by

$$F_{n+2} = F_{n+1} + F_n \qquad F_1 = F_2 = 1$$

```
Solution

fibo <- function(n) {
   res <- integer(n)
   res[1:2] <- 1
   for (k in 3:n) {
      res[k] <- res[k-1] + res[k-2]
   }
   return(res)
}</pre>
fibo(5)

[1] 1 1 2 3 5
```

i Question

Write a function that perform binary search in a non-decreasing vector.

Solution

! 4

Read Chapter on functions in Advanced R

In R, argument evaluation is surprising, powerful but taming argument evaluation is real work.

Functional programming

In R, functions are first class entities, they can be defined at run-time, they can be used as function arguments. You can define list of functions, and iterate over them.

Try to use https://purrr.tidyverse.org.

! Anonymous functions

```
\(x) body
is a shorthand for
function (x) {
  body
}
```

Operators purrr::map_???

i Question

```
Write truth tables for &, |, &&, ||, ! and xor Hint: use purrr::map, function outer()
```

Solution

```
vals <- c(TRUE, FALSE, NA)</pre>
ops <- c(`&`, `|`, `xor`)
truth <- purrr::map(ops, \(x) outer(vals, vals, x))</pre>
names(truth) <- (ops)</pre>
truth
$`.Primitive("&")`
      [,1] [,2] [,3]
[1,] TRUE FALSE
[2,] FALSE FALSE FALSE
[3,]
       NA FALSE
$`.Primitive("|")`
     [,1] [,2] [,3]
[1,] TRUE TRUE TRUE
[2,] TRUE FALSE
[3,] TRUE
          NA
                  NA
\pi(x, y) \n{\ (x | y) & !(x & y)\n}
      [,1] [,2] [,3]
[1,] FALSE TRUE
[2,] TRUE FALSE
                   NA
[3,]
        NA
              NA
                   NA
```

Question

Write a function that takes as input a square matrix and returns TRUE if it is lower triangular.

```
lt <- function(A){</pre>
  n \leftarrow nrow(A)
  all(purrr::map_lgl(1:(n-1), \(x) all(0== A[x, (.+1):n])))
```

Question

Use map, choose and proper use of pronouns to deliver the n first lines of the Pascal triangle using one line of code.

As far as the total number of operations is concerned, would you recommend this way of computing the Pascal triangle?

Solution

```
n <- 5
tp5 <- matrix(unlist(map(0:n,</pre>
           (x) c(choose(x, 0:x), rep(0L, n-x)))),
       nrow=n+1,
       byrow=T)
rownames(tp5) <- 0:n
colnames(tp5) <- 0:n
tp5
  0 1
       2 3 4 5
0 1 0 0 0 0 0
1 1 1 0 0 0 0
2 1 2 1 0 0 0
3 1 3 3 1 0 0
4 1 4 6 4 1 0
5 1 5 10 10 5 1
No. Using map and choose, we do not reuse previous computations. The total number
```

of arithmetic operations is $\Omega(n^3)$, it should be $O(n^2)$.

Read Chapter on Functional Programming in Advanced R

Further exploration

This notebook walked you through some aspects of R and its packages. We just saw the tip of the iceberg.

We barely mentioned:

- (Non-standard) Lazy evaluation
- Different flavors of object oriented programming
- Connection with C++: RCpp
- Connection with databases: dbplyr
- Building modeling pipelines: tidymodels
- Concurrency

- Building packages
- Building interactive Apps: Shiny
- Attributes (metadata)
- Formulae formula
- Strings stringi, stringr
- Dates lubridate
- and plenty other things

References

- https://www.statmethods.net/index.html
- https://www.datacamp.com/courses/free-introduction-to-r
- dplyr videos
- ggplot2 video tutorial
- cheatsheets
- A Readers who really want to learn R should spend time on
 - R for Data Science by Wickham, Cetinkaya-Rundel, and Grolemund.
 - Advanced R 2nd Edition by Wickham
 - Advanced R Solutions by Grosser and Bumann
 - Hands-On Programming with R by Grolemund

Don't go without Base R cheatsheet